

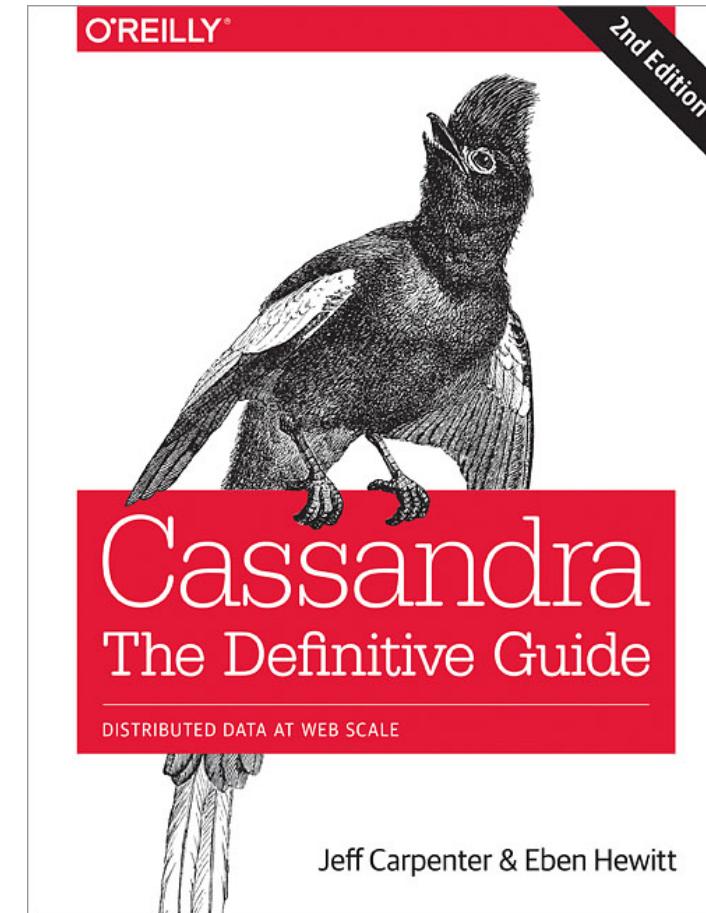
# Building Applications with Apache Cassandra

A quickstart guide to Cassandra for Java developers and architects



# Hoping You Came Prepared...

- Download Virtual Machine for Activities
  - If not, please take a minute now and start the download
  - See Slack/chat for URL
- Review Cassandra, The Definitive Guide, 2nd Ed
  - Chapter 1 – Beyond Relational Databases
  - Chapter 2 – Introducing Cassandra
  - Chapter 4 – The Cassandra Query Language (through “Cassandra’s Data Model”)



# Goals

**By the end of this live, online course, you'll understand:**

- Cassandra's architecture and how it works
- How Cassandra is different from traditional databases, and how those differences affect data modeling and application development
- How to identify use cases where Cassandra is a good fit
- Cassandra's features, especially as available through the client drivers
- Commonly used configuration options

# Goals

**By the end of this live, online course, you'll be able to:**

- Design Cassandra data models that will perform / scale effectively
- Create application programs using the DataStax Java Driver
- Use tools including *cqlsh*, *nodetool*, and CCM

# Non-Goals

- Advanced configuration options (see Chapter 7)
- Monitoring beyond *nodetool* basics (see Chapter 10)
- Maintaining Cassandra clusters (see Chapter 11)
- Performance tuning (see chapter 12)
- Security beyond basic client authentication (see Chapter 13)
- Integrating Cassandra with other technologies (see chapter 14 and [datastax.com](http://datastax.com))



# Who am I?

- Developer
- Architect
- Author
- Developer Advocate
- Defense
- Hospitality
- R&D
- Distributed Systems
- Large Scale
- Cassandra

# Contact Info

[jeff.carpenter@datastax.com](mailto:jeff.carpenter@datastax.com)

 @jscarp

 jeffreyscarpenter

Blog: [medium.com/@jscarp](https://medium.com/@jscarp)



# Unit 1: Introduction, Install and *cqlsh*

~50 minutes

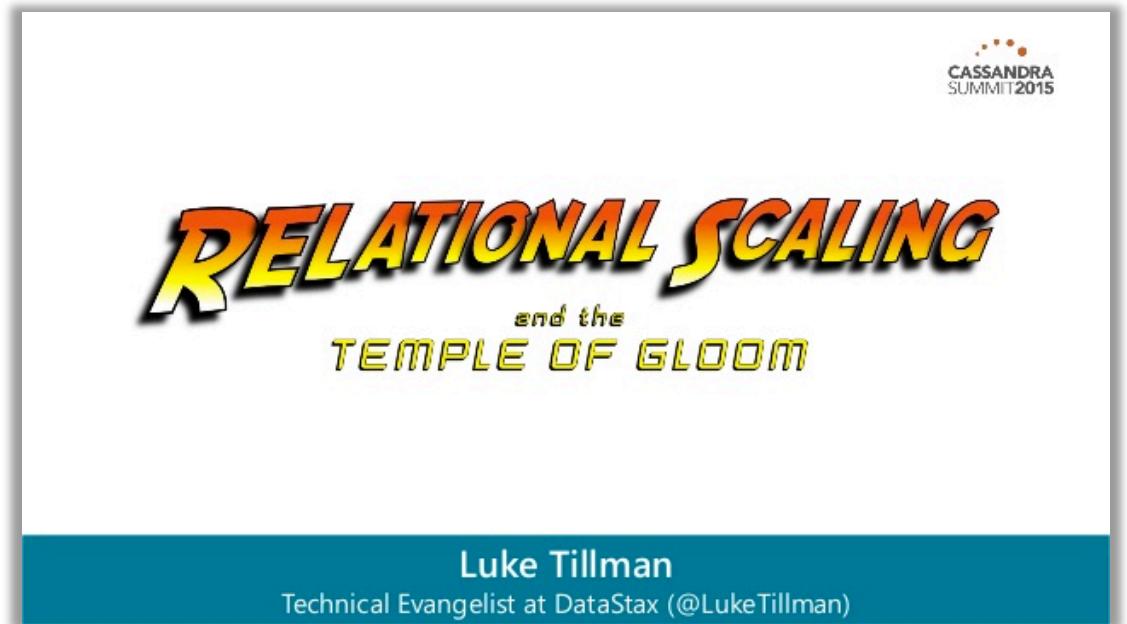
# Unit 1 Goals

- Understand the challenges that come with scaling traditional RDBMS
- Articulate how Cassandra is different and the use cases for which it is a good fit
- Get a good “hands on” experience with installing Cassandra and running basic commands in *cqlsh*

# What's wrong with RDBMS?

- Pros
  - Relational data modeling is well understood
  - SQL is easy to use and ubiquitous
  - ACID transactions - data integrity
- Cons
  - Scaling is hard, sharding and replication have side-effects (performance, reliability, cost)
  - Trend toward denormalization

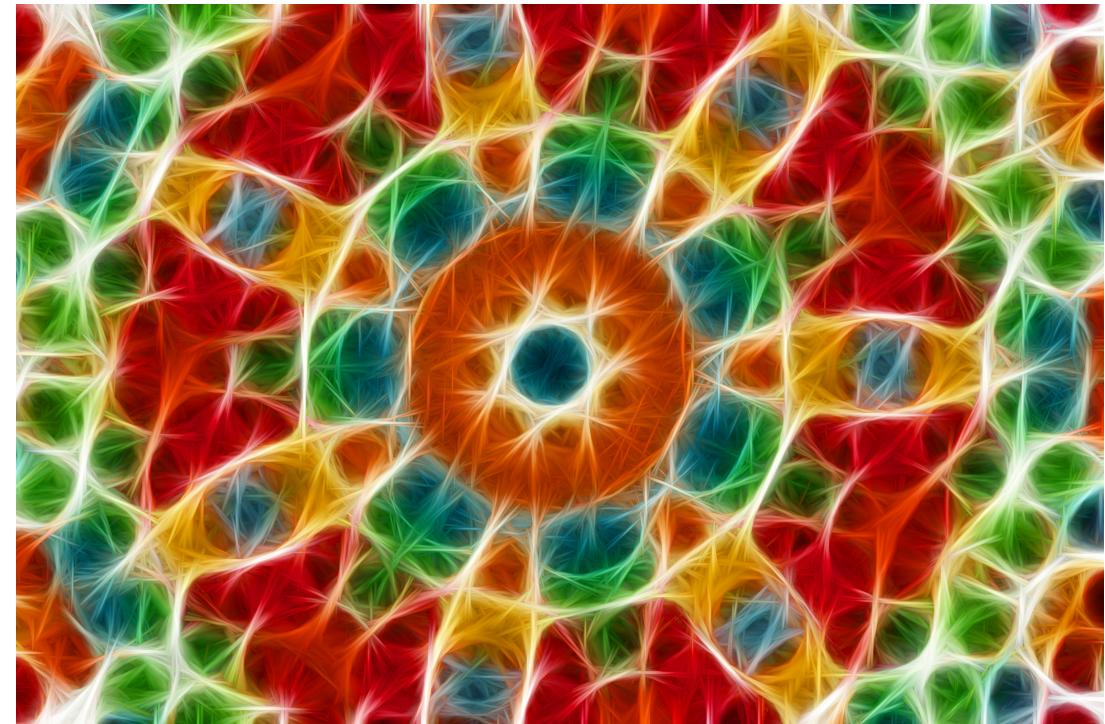
For fun: [Relational Scaling and the Temple of Gloom](#)



See also: Cassandra, The Definitive Guide, 2<sup>nd</sup> Edition, Chapter 1: Beyond Relational Databases

# The NoSQL Revolution (~2009)

- Key-value – Dynamo, Riak, Voldemort, Redis, Memcached
- Column-oriented – BigTable, Hbase, Cassandra\*
- Document – Mongo DB, Document DB, CouchDB
- Graph – Neo4J, DSE Graph
- Multi-model – DataStax Enterprise, CosmosDB
- Relational renaissance – AWS Aurora, Google Spanner



See also: [How Many Databases?](#)

# Apache Cassandra at a Glance

- First developed by Facebook
- Became a top-level Apache Foundation project in 2010
- Distributed, decentralized
- Elastic scalability / high performance
- High availability / fault tolerant
- Tuneable consistency
- Partitioned row store



See also: Cassandra, The Definitive Guide, 2<sup>nd</sup> Edition, Chapter 2: Introducing Cassandra

# Problems Cassandra is Especially Good At

- Large scale storage
  - >10s of TB
- Lots of writes
  - Time-series data, IoT
- Statistics and analytics
  - For example, as a Spark data source
- Geographic distribution
  - Multiple data centers



# Discussion: Is Cassandra a Good fit for my project?

- What region are you from?
- What industry are you in?
- What are your use cases?
- Feel free to share in the chat (briefly!)



# Exercise 1a - Installing and Starting Cassandra

- Goal:
  - Gain familiarity with the contents of the Apache Cassandra distribution
- Steps:
  - Start VM
  - Cassandra archive is located in the Downloads folder
  - Open a terminal and expand the archive in your home directory:

```
$ tar xvf Downloads/apache-cassandra-3.11.3-bin.tar.gz
```

- Examine the directory

```
$ cd apache-cassandra-3.11.3
```

# Exercise 1a - Installing and Starting Cassandra

- Examining the Installation:
  - Note the contents of the following directories, and then we'll discuss:
    - /bin
    - /conf
    - /doc
    - /lib
    - /tools
  - See if you can find the answers to these questions:
    - How do I start Cassandra?
    - Where can I find a list of what is fixed in this release?
    - Which config files are required to be present?
    - How would I learn about the Cassandra Query Language (CQL) syntax?
    - Where should I go to learn how to troubleshoot my installation?

# Examining the Installation - Highlights

- **/bin**
  - Scripts to start/stop Cassandra, nodetool, utilities for maintaining SSTables
- **/conf**
  - Configuration files: `cassandra.yaml`,
- **/doc**
  - Quick reference documentation for Cassandra and CQL
- **/lib**
  - Executables for Cassandra and libraries on which it depends
- **/tools**
  - Tools for stress testing and analyzing SSTables

What's missing?

# Exercise 1a - Installing and Starting Cassandra

- Questions and Answers
  - How do I start Cassandra?
    - /bin/cassandra
  - Where can I find a list of what is fixed in this release?
    - CHANGES.txt
  - Which config files are required to be present?
    - /conf/cassandra.yaml, logback.xml
  - How would I learn about the Cassandra Query Language (CQL) syntax?
    - /doc/cql3/CQL.html
  - Where should I go to learn how to troubleshoot my installation?
    - /doc/html/troubleshooting/index.html

# Exercise 1a - Installing and Starting Cassandra

- Starting Cassandra:
  - In a terminal window, start Cassandra

```
$ bin/cassandra -f
```

- This will cause Cassandra to run in the foreground so you can see its logs
- Not a normal production setting, but useful for demonstration

# Getting Cassandra

- Distributions
  - Apache Cassandra
  - DataStax Enterprise
- Local Installation
  - Manual
  - Build from source (Chapter 3)
  - Apt-get (Ubuntu)
  - Homebrew (Mac)
- Deployment Options
  - Docker images
  - Virtual Machines
  - Mesosphere DC/OS (Open source and DSE)
  - Cloud marketplaces (AWS, Azure, Google Cloud)
  - Managed services (DataStax Managed Cloud, Instaclustr)

# Exercise 1b – Running *cqlsh*

- Goal:
  - Learn how to use *cqlsh* including commonly used commands
- Steps
  - In a terminal window, start *cqlsh*

```
$ bin/cqlsh
```

- The shell automatically connects to localhost at default port (9042)
- Use "HELP" command will list available commands
- For a list of sample commands to try out, see this file

```
~/cassandra-guide/resources/cqlsh_intro.cql
```

See also: <https://github.com/jeffreyscarpenter/cassandra-guide>

# CQL Commands

- Data Definition (DDL)
  - CREATE KEYSPACE
  - ALTER KEYSPACE
  - DROP KEYSPACE
  - CREATE TABLE
  - ALTER TABLE
  - DROP TABLE
  - TRUNCATE
  - USE
- Data Manipulation (DML)
  - SELECT
  - INSERT
  - UPDATE
  - DELETE
  - BATCH

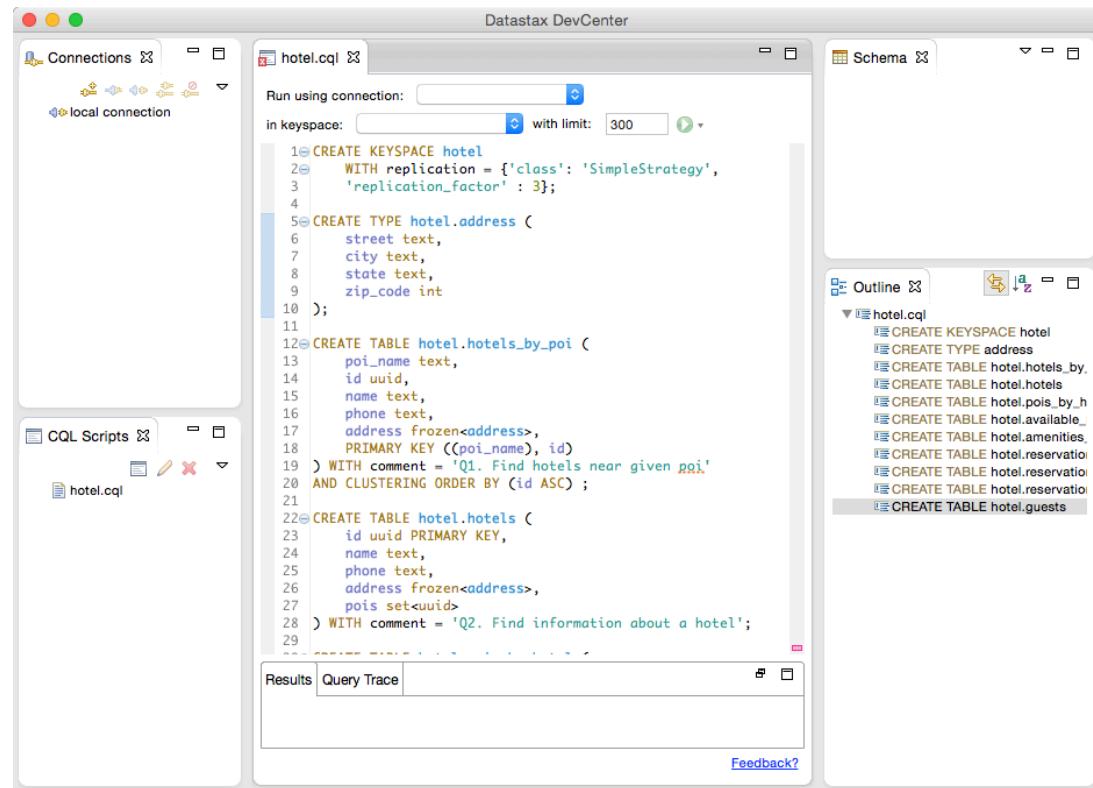
# CQL Commands (continued)

- Indexes and Materialized Views
  - CREATE INDEX
  - DROP INDEX
  - CREATE MATERIALIZED VIEW
  - ALTER MATERIALIZED VIEW
  - DROP MATERIALIZED VIEW
- Security
  - CREATE USER
  - ALTER USER
  - DROP USER
  - CREATE ROLE
  - ALTER ROLE
  - DROP ROLE
  - GRANT
  - REVOKE
  - LIST PERMISSIONS
  - LIST ROLES
  - LIST USERS

# *Cqlsh*-Only Commands

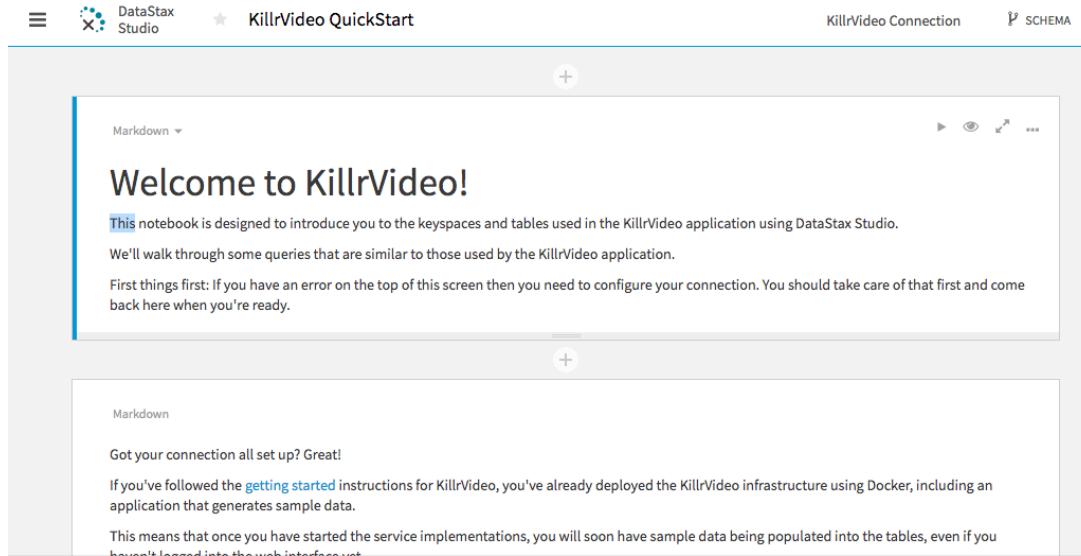
- Default settings
  - CONSISTENCY
  - SERIAL CONSISTENCY
  - PAGING
  - TRACING
  - EXPAND
- Metadata
  - DESCRIBE KEYSPACE
  - DESCRIBE TABLE
- File manipulation
  - COPY
  - SOURCE
  - CAPTURE
- Other commands
  - CLEAR
  - HELP
  - EXIT

# DataStax DevCenter



- Graphical, Eclipse-style tool for editing schema and executing CQL
- Works with Apache Cassandra and DataStax Enterprise
- No longer supported, superseded by DataStax Studio

# DataStax Studio



- Web application, notebook-style tool for editing schema and executing CQL
- Also supports Graph (Gremlin) schema and queries
- Works with DataStax Enterprise only

# Break

5 minutes

# Unit 2: Data Modeling

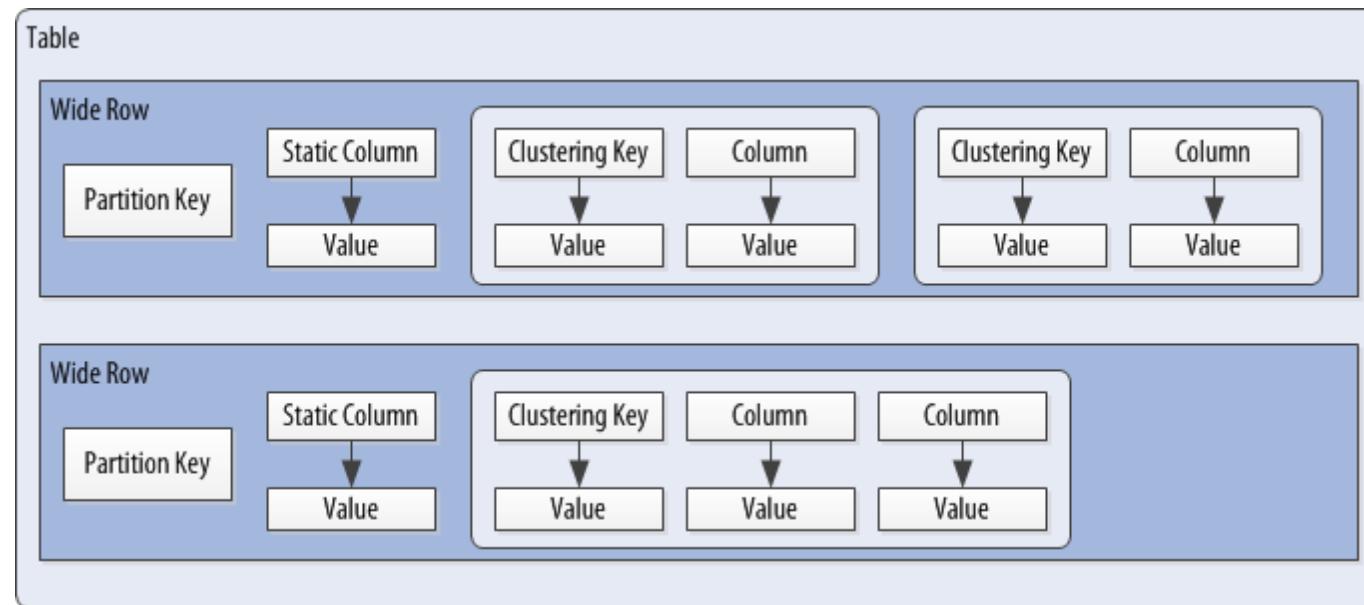
~55 minutes

# Unit 2 Goals

- Familiarity with Cassandra's data model and terminology
- Learn what data types are built into Cassandra, how they map to Java types, and how you can extend Cassandra with your own data types
- Understand the basic principles of Cassandra data modeling and how it is different from traditional modeling
- Get practical experience designing data models for real applications using Cassandra

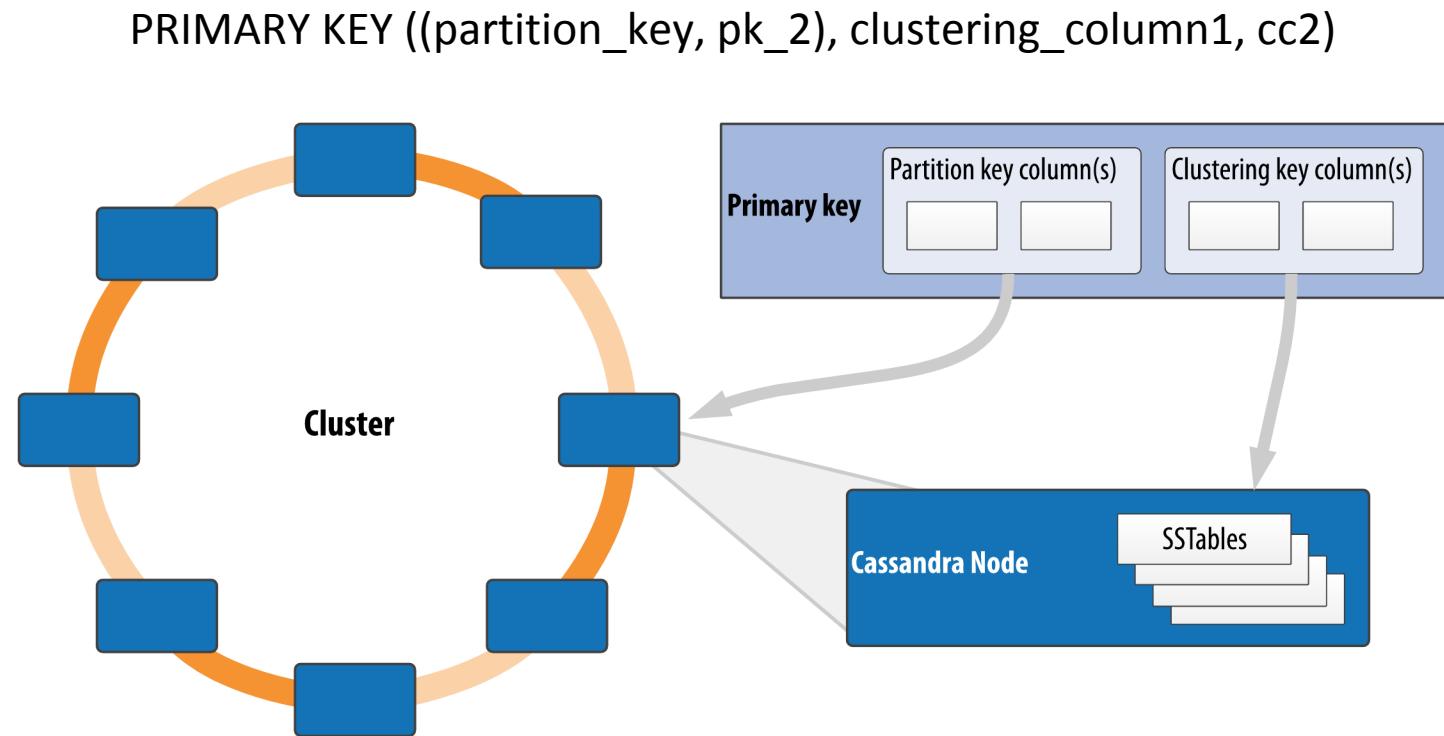
# Terminology - Cassandra's Data Model

A Cassandra table



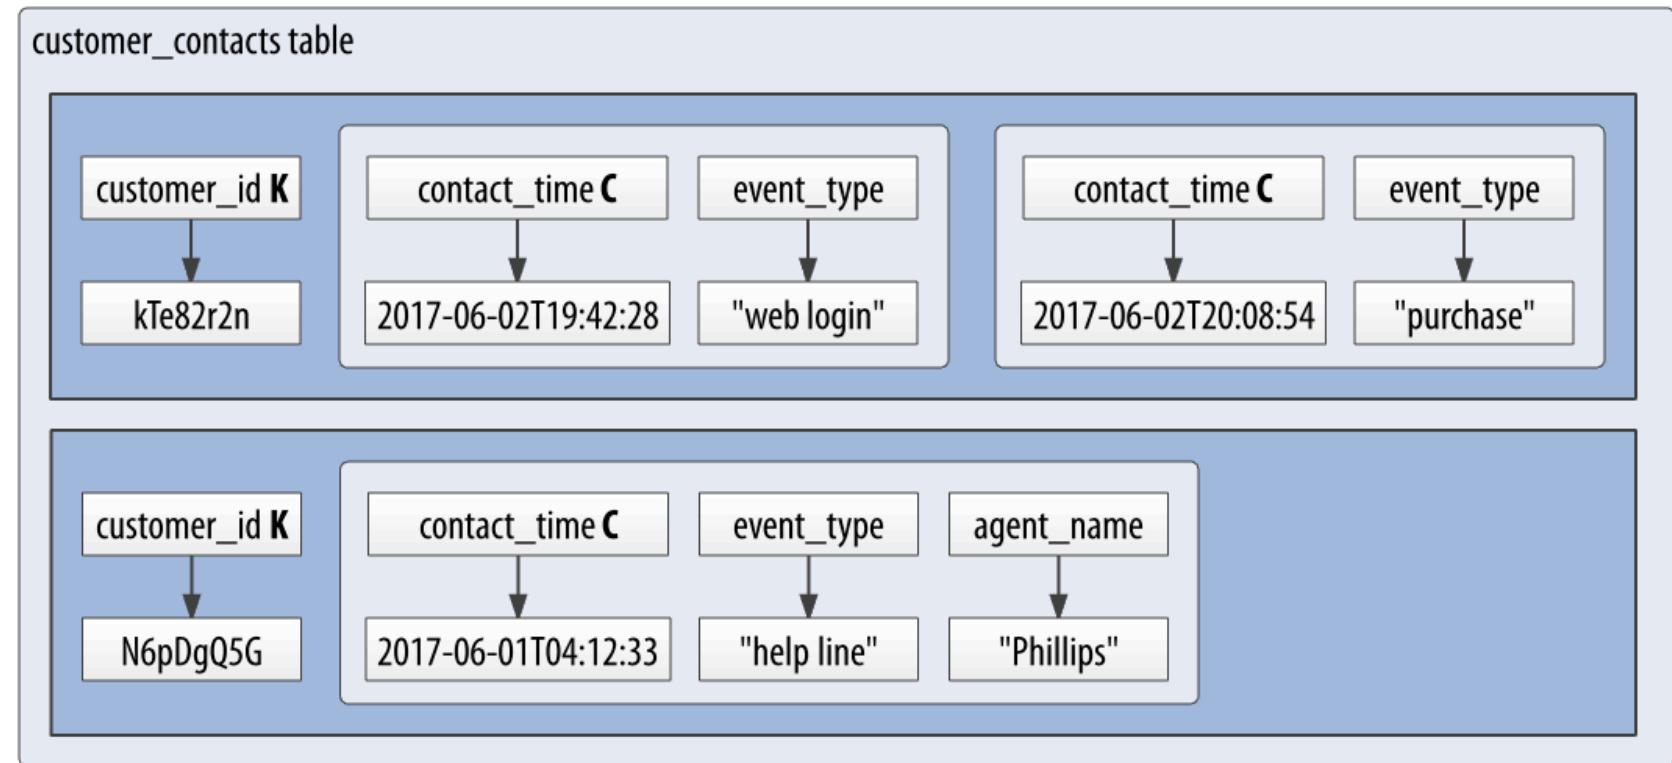
See also: Cassandra, The Definitive Guide, 2<sup>nd</sup> Edition, Chapter 4: The Cassandra Query Language

# Terminology – Partition and Clustering Keys



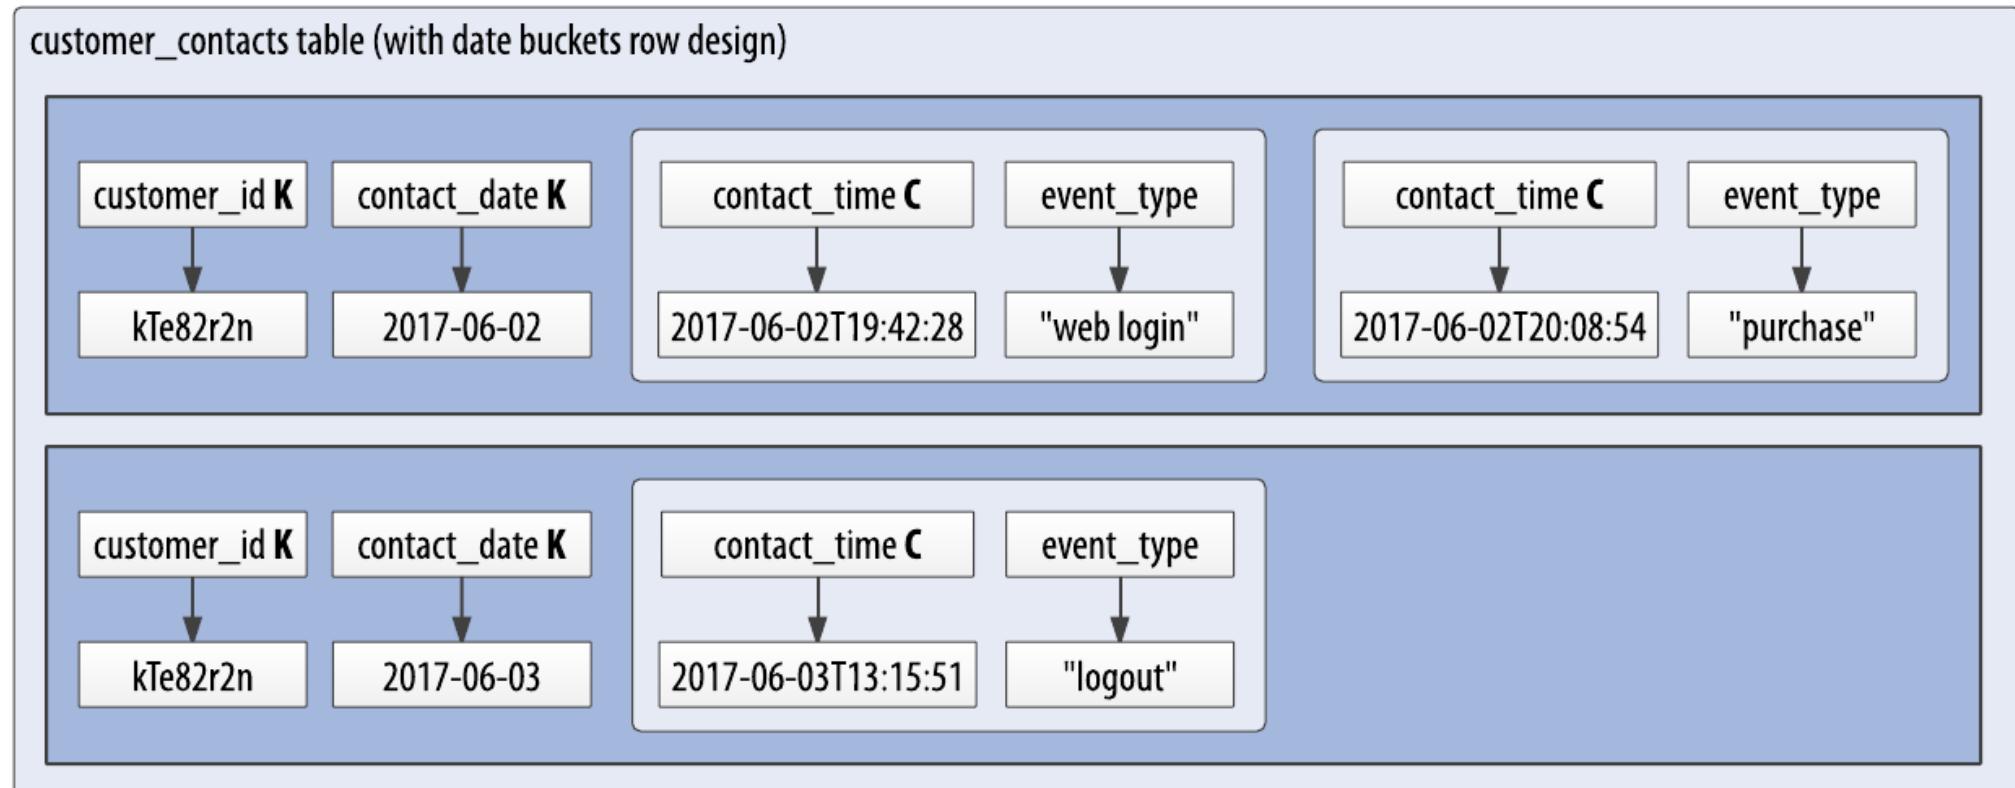
# Partition and Clustering Key Example

Labeling:  
**K** – partition key  
**C** – clustering key



PRIMARY KEY ((customer\_id), contact\_time)

# Buckets – a way to avoid overly wide rows



PRIMARY KEY ((**customer\_id, contact\_date**), contact\_time)

# Exercise 2a – Cassandra Data Types

- Goal:
  - Create and manipulate tables using *cqlsh*
  - Learn how to use Cassandra’s built-in types and collections
- Steps
  - Make sure *cqlsh* is running
  - Work through the commands in the file

```
~/cassandra-guide/resources/learning_cql_data_types.cql
```

# Cassandra Data Types

## Numeric Data Types

CQL Type	Description	Java Type
int	32-bit signed integer	int
bigint	64-bit signed integer	long
smallint	16-bit signed integer	short
tinyint	8-bit signed integer	tinyint
varint	Variable-precision signed integer	java.math.BigInteger
float	32-bit floating point	float
double	64-bit floating point	double
decimal	Variable-precision decimal	java.math.BigDecimal

## Text Data Types

CQL Type	Description	Java Type
text, varchar	UTF-8 character string	java.lang.String
ascii	ASCII character string	java.lang.String

See also:

<http://docs.datastax.com/en/developer/java-driver/3.6/manual/>

# Cassandra Data Types

## Time / Identity types

CQL Type	Description	Java Type
timestamp	ISO 8601 date/time, encoded as 64-bit int	java.util.Date
date	ISO 8601 date only	java.util.LocalDate
time	ISO 8601 time only	long
uuid	Type 4 Universal Unique Identifier	java.util.UUID
timeuuid	Type 1 UUID	java.util.UUID

## Miscellaneous types

CQL Type	Description	Java Type
boolean	True/false value	boolean
blob	Arbitrary array of bytes	java.nio.ByteBuffer
inet	IP v4/v6 address	java.net.InetAddress
counter	64-bit integer which can only be incremented or decremented	long

# Cassandra Data Types

## Collections

CQL Type	Description	Java Type
set	Unordered collection of elements	java.util.Set

```
// set creation
ALTER TABLE user ADD emails set<text>
```

```
// populate set
UPDATE user SET emails = {
    'mary@example.com' } WHERE ...

// add element
UPDATE user SET emails = emails + {
    'mary@example.com' } WHERE ...

// Remove element
UPDATE user SET emails = emails - {
    'mary@example.com' } WHERE ...
```

See also: [http://docs.datastax.com/en/cql/3.3/cql/cql\\_reference/collection\\_type\\_r.html](http://docs.datastax.com/en/cql/3.3/cql/cql_reference/collection_type_r.html)

# Cassandra Data Types

## Collections

CQL Type	Description	Java Type
list	Ordered collection of elements	java.util.List

```
// list creation
ALTER TABLE user ADD phone_numbers
list<text>
```

```
// populate list
UPDATE user SET phone_numbers = [
    '1-800-999-999' ] WHERE ...

// append element
UPDATE user SET emails = phone_numbers + {
    '480-441-6957' } WHERE ...

// set item
UPDATE user SET emails[0] = '480-111-1111'
WHERE ...
```

See also: [http://docs.datastax.com/en/cql/3.3/cql/cql\\_reference/collection\\_type\\_r.html](http://docs.datastax.com/en/cql/3.3/cql/cql_reference/collection_type_r.html)

# Cassandra Data Types

## Collections

CQL Type	Description	Java Type
map	Collection of key/value pairs	java.util.Map

```
// map creation
ALTER TABLE user ADD login_sessions
map<timeuuid, int>
```

```
// populate map
UPDATE user SET login_sessions = {
  6061b850..., 1 } WHERE ...

// add element
UPDATE user SET login_sessions =
  login_sessions + { a9fac1d0..., 2} WHERE ...

// Remove element
UPDATE user SET login_sessions =
  login_sessions - { a9fac1d0... } WHERE ...
```

See also: [http://docs.datastax.com/en/cql/3.3/cql/cql\\_reference/collection\\_type\\_r.html](http://docs.datastax.com/en/cql/3.3/cql/cql_reference/collection_type_r.html)

# Cassandra Data Types

## Tuple

CQL Type	Description	Java Type
tuple	Fixed structure of other data types	com.datastax.driver.core.TupleType

```
// tuple creation
ALTER TABLE user ADD address
<tuple<text, text, text, text, text>>;
```

```
// populate tuple
UPDATE user SET address = (
    '7712 E Broadway', // street address
    'Tucson',          // city
    'AZ',              // state or province
    '85715',           // postal code
    'USA'              // country
) WHERE ...
```

See also: [http://docs.datastax.com/en/cql/3.3/cql/cql\\_reference/tupleType.html](http://docs.datastax.com/en/cql/3.3/cql/cql_reference/tupleType.html)

# User Defined Types

- Scoped to keyspace
- Can be nested
- Similar syntax to defining table

```
CREATE TYPE address (
    street text,
    city text,
    state_or_province text,
    postal_code text,
    country text
);
```

```
ALTER TABLE user ADD address address;

UPDATE user SET address =
    street: '7712 E Broadway',
    city: 'Tucson',
    state_or_providence: 'AZ',
    postal_code: '85715',
    country: 'USA'
) WHERE ...
```

# Exercise 2b - Query Options

- Goal:
  - Learn about the various clauses available in the CQL SELECT statement (WHERE, IN, SORT)
- Steps
  - Make sure *cqlsh* is running
  - Work through the commands in the file

```
~/cassandra-guide/resources/learning_cql_query_clauses.cql
```

# JSON Input / Output

- Convenient way to insert/retrieve data
- Can save processing JSON input in your application
- Not "schemaless"!

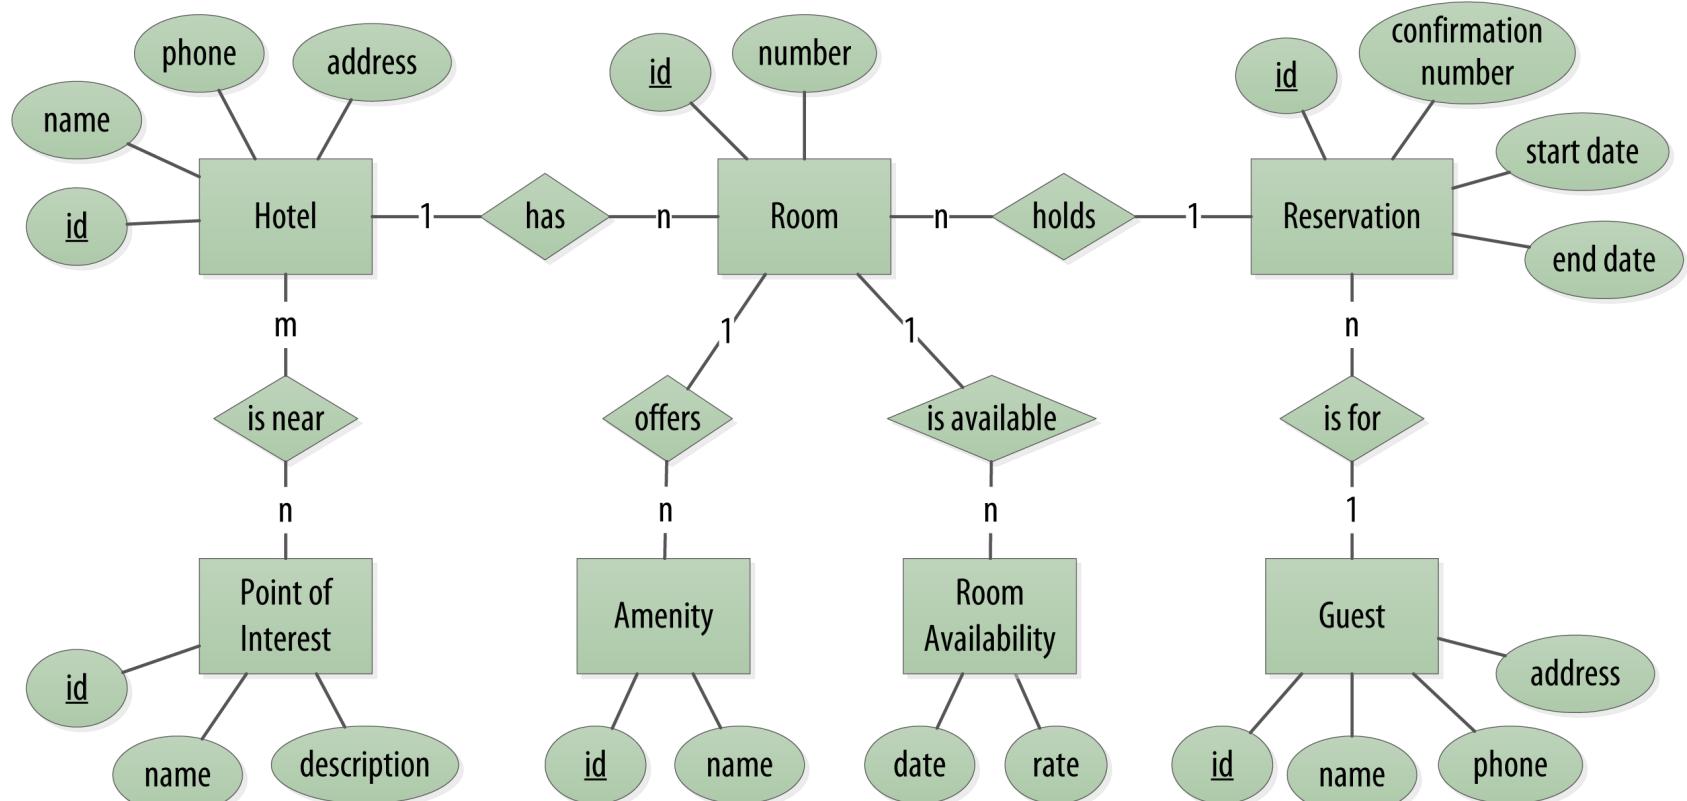
```
INSERT INTO user JSON '{"first_name": "Bill",  
"last_name": "Nguyen", : "title": "Mr."}';  
  
SELECT JSON * FROM user;  
[json]  
-----  
{"first_name": "Mary", "last_name":  
"Rodriguez", : "title": null}  
 {"first_name": "Bill", "last_name":  
"Nguyen", : "title": "Mr."}
```

# Cassandra Data Modeling

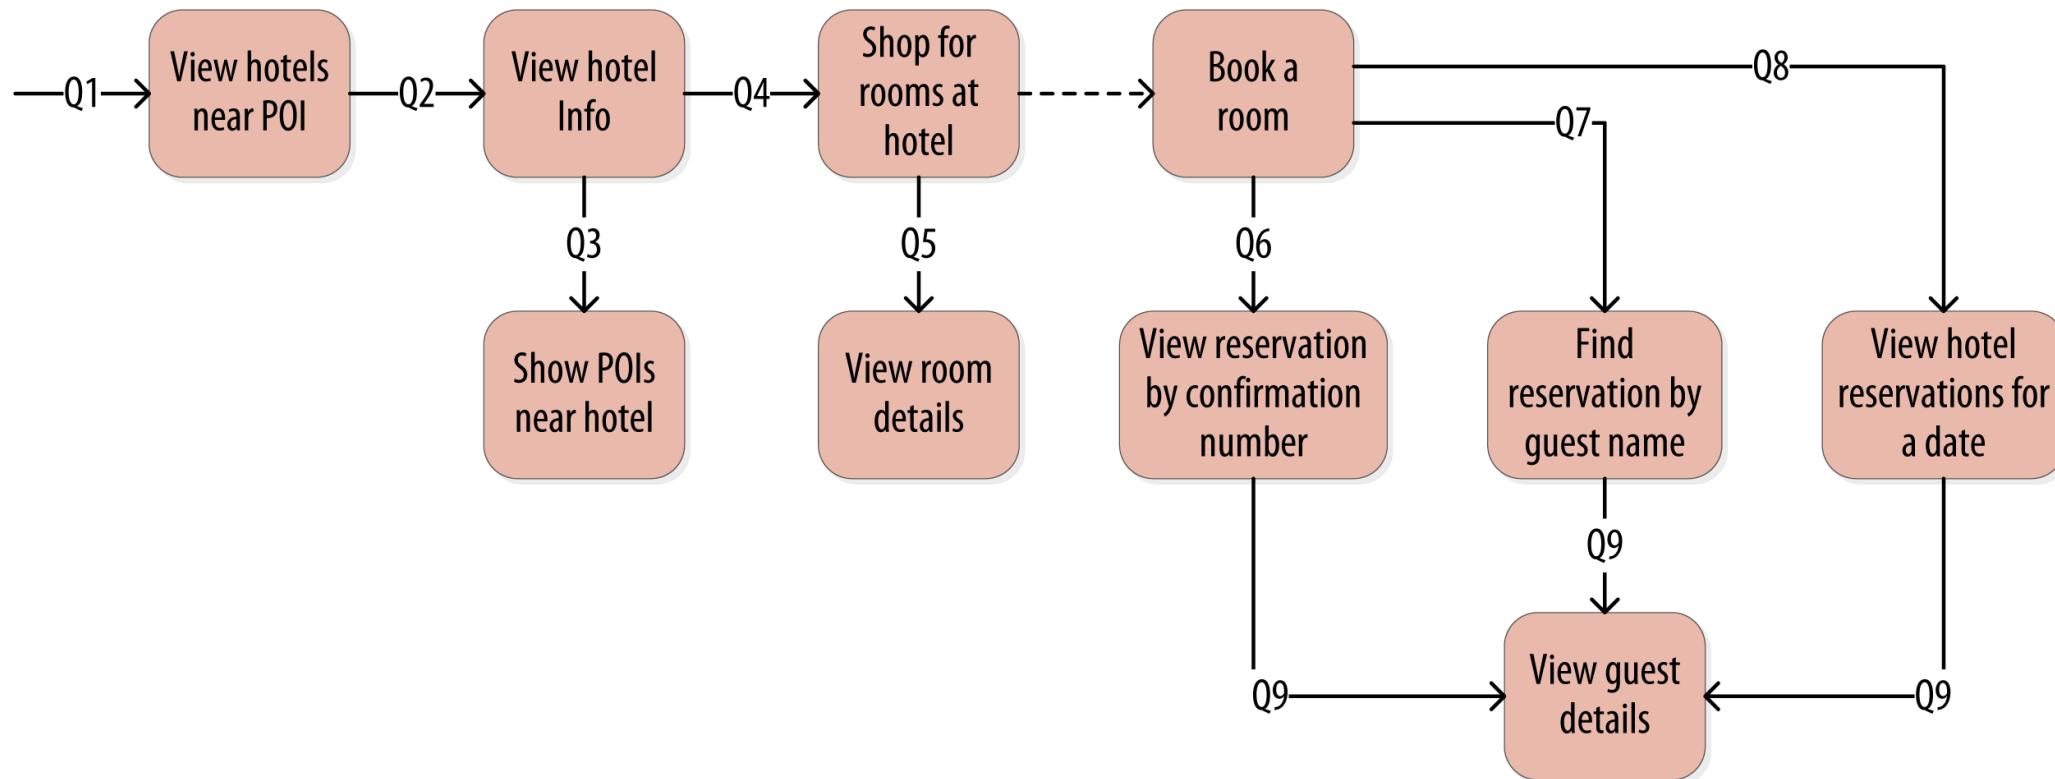
- Classic data modeling approach
  - Conceptual
  - Logical
  - Physical
- Additional Cassandra data modeling steps
  - Identifying access patterns
  - Evaluating data models

See also: [Cassandra, The Definitive Guide, 2<sup>nd</sup> Edition, Chapter 5: Data Modeling](#)

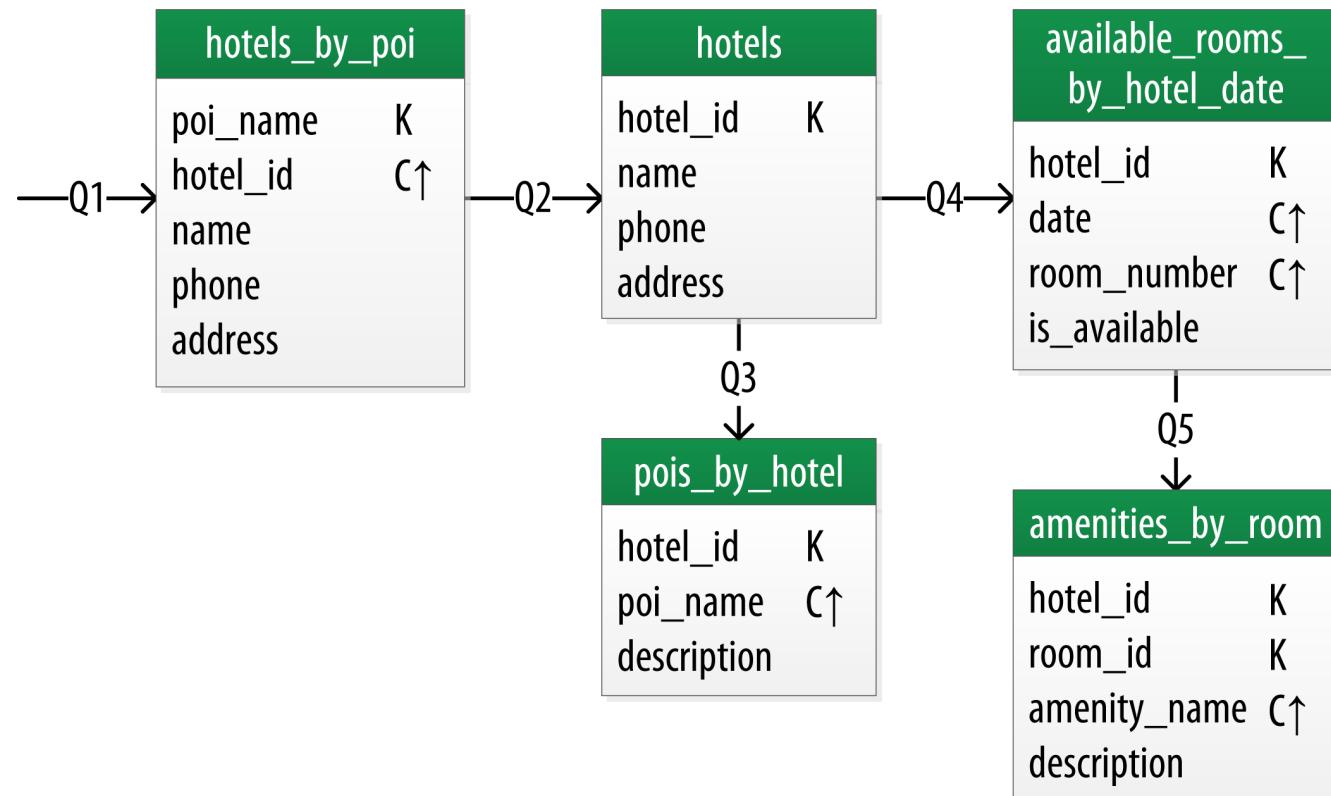
# Conceptual Data Model – Hotel Domain



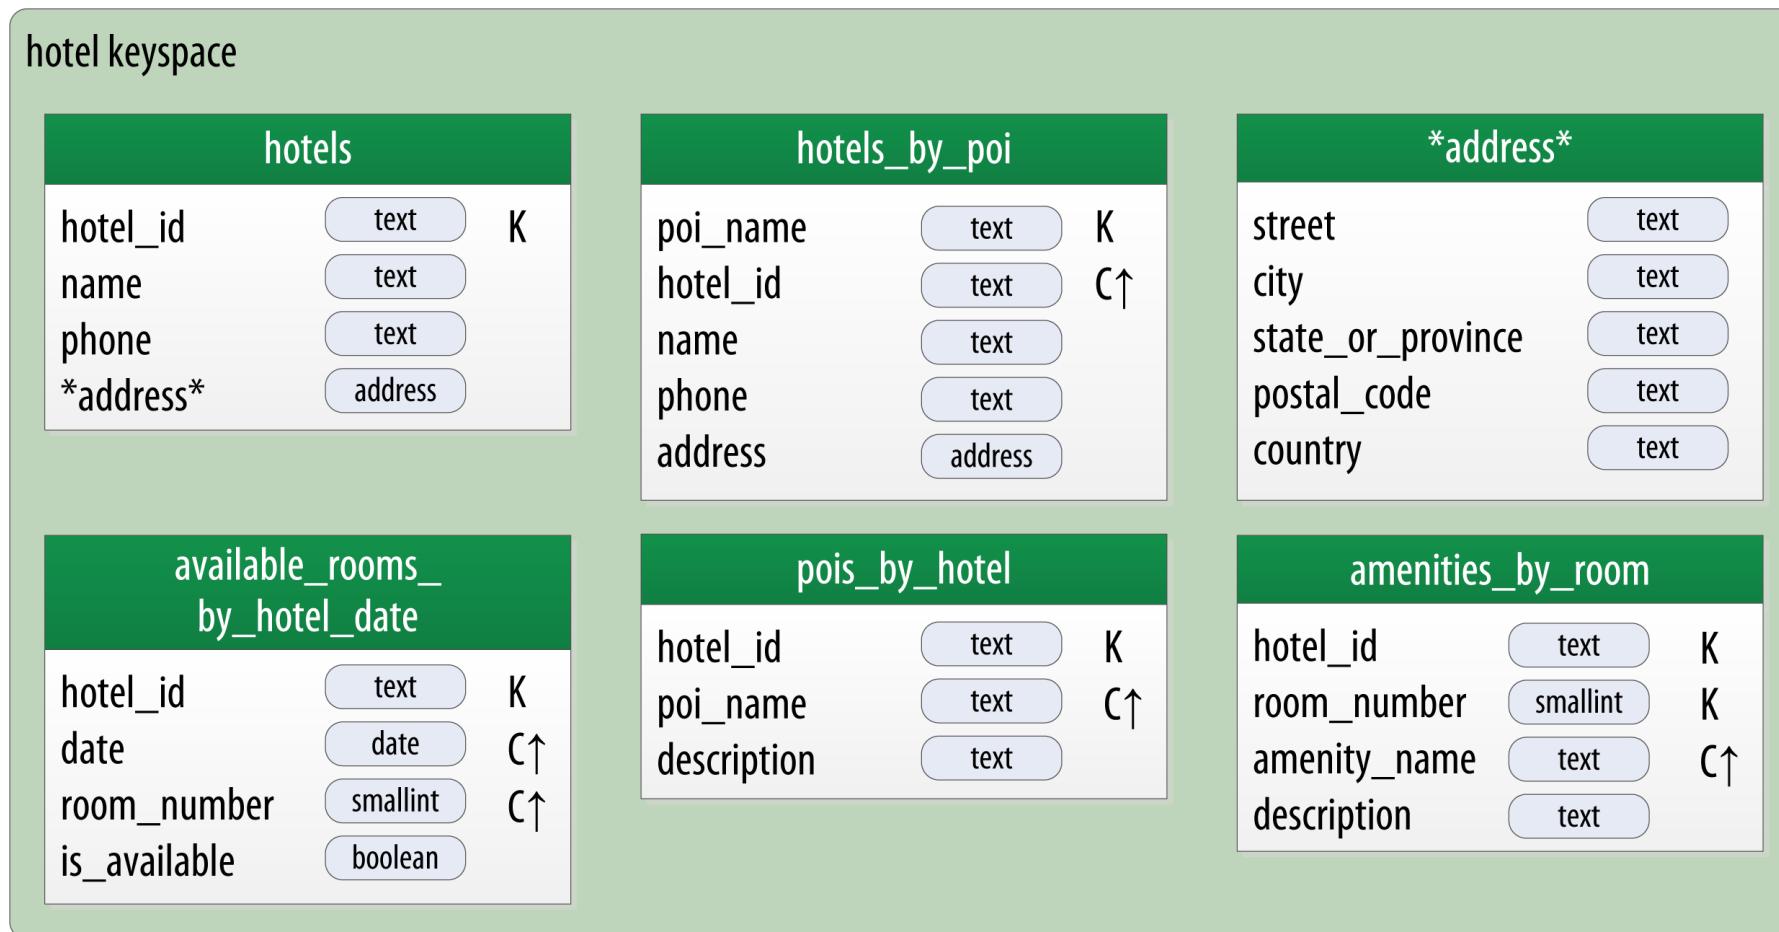
# Access Patterns – Hotel Domain



# Logical Data Modeling – Chebotko Diagram



# Physical Data Model – Chebotko Diagram



# CQL Schemas

```
CREATE KEYSPACE hotel  
WITH replication = {'class':  
'SimpleStrategy',  
'replication_factor' : 3};
```

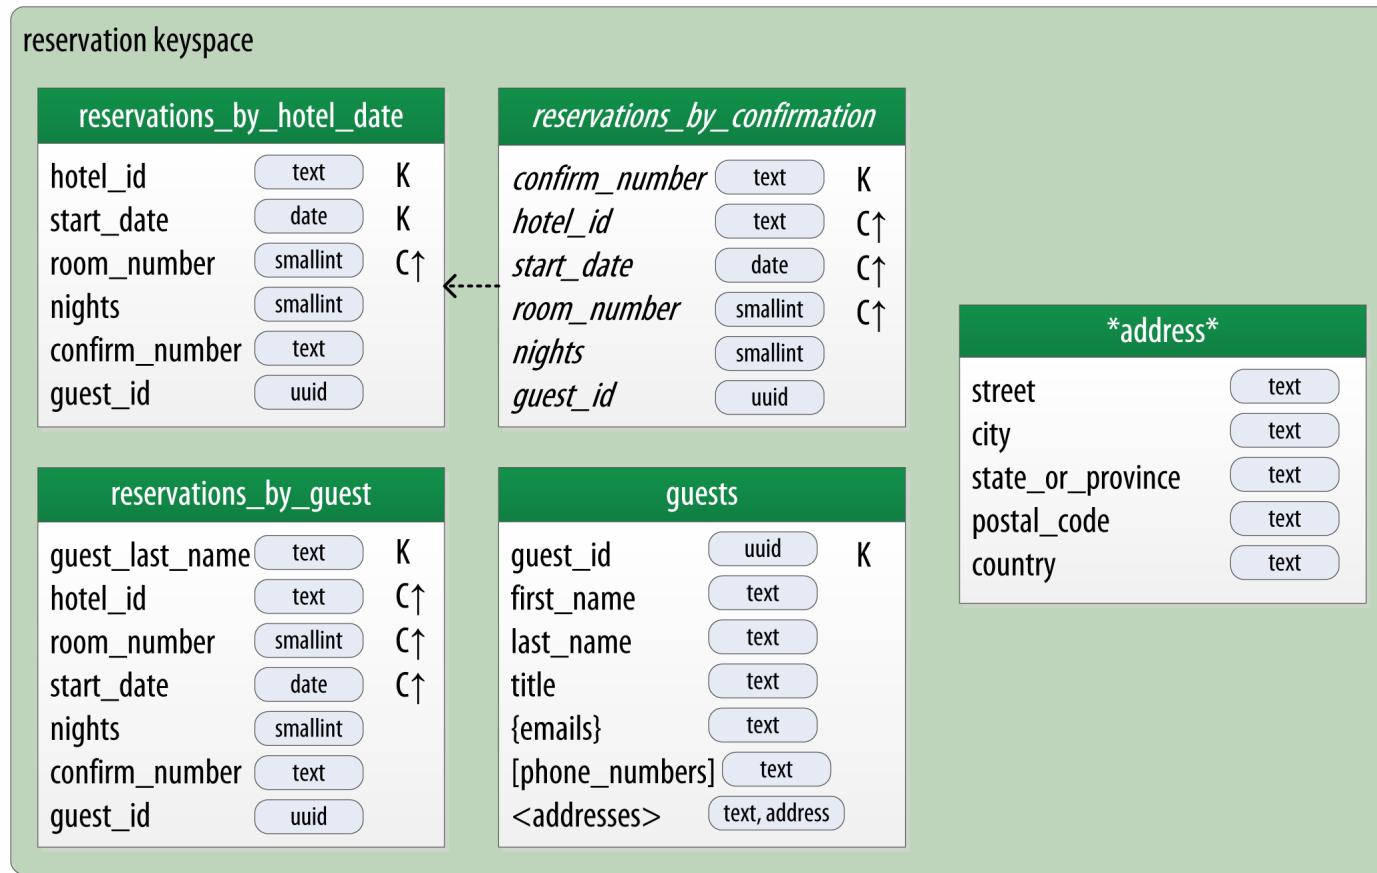
```
CREATE TYPE hotel.address (  
    street text,  
    city text,  
    state_or_province text,  
    postal_code text,  
    country text);
```

```
CREATE TABLE hotel.hotels_by_poi (  
    poi_name text,  
    hotel_id text,  
    name text,  
    phone text,  
    address frozen<address>,  
    PRIMARY KEY ((poi_name),  
                 hotel_id)  
)  
WITH CLUSTERING ORDER BY (  
    hotel_id ASC) ;
```

# Exercise 2c - Data Modeling

- Goal:
  - Create a schema for hotel reservations based on the design provided
- Steps
  - Create a keyspace for reservations using the SimpleReplication strategy and a replication factor of 1
  - Define a UDT for the Address data type
  - Define tables for reservations\_by\_hotel\_date and reservations\_by\_guest
  - Create reservations\_by\_confirmation as a materialized view on reservations\_by\_hotel\_guest

# Exercise 2c - Data Modeling



# Data Modeling Challenges

- What were some of the struggles you encountered with creating the data model (not tool related)?
- If you did have any tool challenges, record them in the chat and I will follow up later on.

# More Cassandra Data Modeling Resources



## Training

Free online training at DataStax Academy



## Reference Application

See a real data model in this working reference application



## Reading

Check out my blog on advanced data modeling: "Data Model Meets World"

See also: *Cassandra, The Definitive Guide, 2<sup>nd</sup> Edition, Chapter 5: Data Modeling*

# Break

5 minutes

# Unit 3: Client Development

~60 minutes

# Unit 3 Goals

- Learn about the various language drivers available for writing Cassandra client applications and the common features they share
- Get “hands on” experience writing, configuring, and running a Java application using the DataStax Java Driver
- Understand the various methods of executing statements supported by the driver and how to use the mapper

# Client Drivers

- Historic Drivers Developed by various community projects, mostly based on Thrift and now deprecated
  - Hector (Java), Astyanax (Java, Netflix), Pycassa (Python), Perlcassa (Perl), Helenus (node.js), Cassandra-Sharp (C#/.NET)
- DataStax Drivers
  - Based on CQL
  - Open Source - <https://github.com/datastax>

See also: Cassandra, The Definitive Guide, 2<sup>nd</sup> Edition, Chapter 8: Clients

# Client Drivers

Apache Cassandra Drivers (Open Source)	DataStax Drivers (DataStax Enterprise)
<a href="#"><u>DataStax Java Driver</u></a>	<a href="#"><u>DataStax Enterprise Java Driver</u></a>
<a href="#"><u>DataStax Python Driver</u></a>	<a href="#"><u>DataStax Enterprise Python Driver</u></a>
<a href="#"><u>DataStax Node.js Driver</u></a>	<a href="#"><u>DataStax Enterprise Node.js Driver</u></a>
<a href="#"><u>DataStax Ruby Driver</u></a>	<a href="#"><u>DataStax Enterprise Ruby Driver</u></a>
<a href="#"><u>DataStax C# Driver</u></a>	<a href="#"><u>DataStax Enterprise C# Driver</u></a>
<a href="#"><u>DataStax C/C++ Driver</u></a>	<a href="#"><u>DataStax Enterprise C/C++ Driver</u></a>
<a href="#"><u>DataStax PHP Driver</u></a>	<a href="#"><u>DataStax Enterprise PHP Driver</u></a>

<https://academy.datastax.com/downloads/download-drivers>

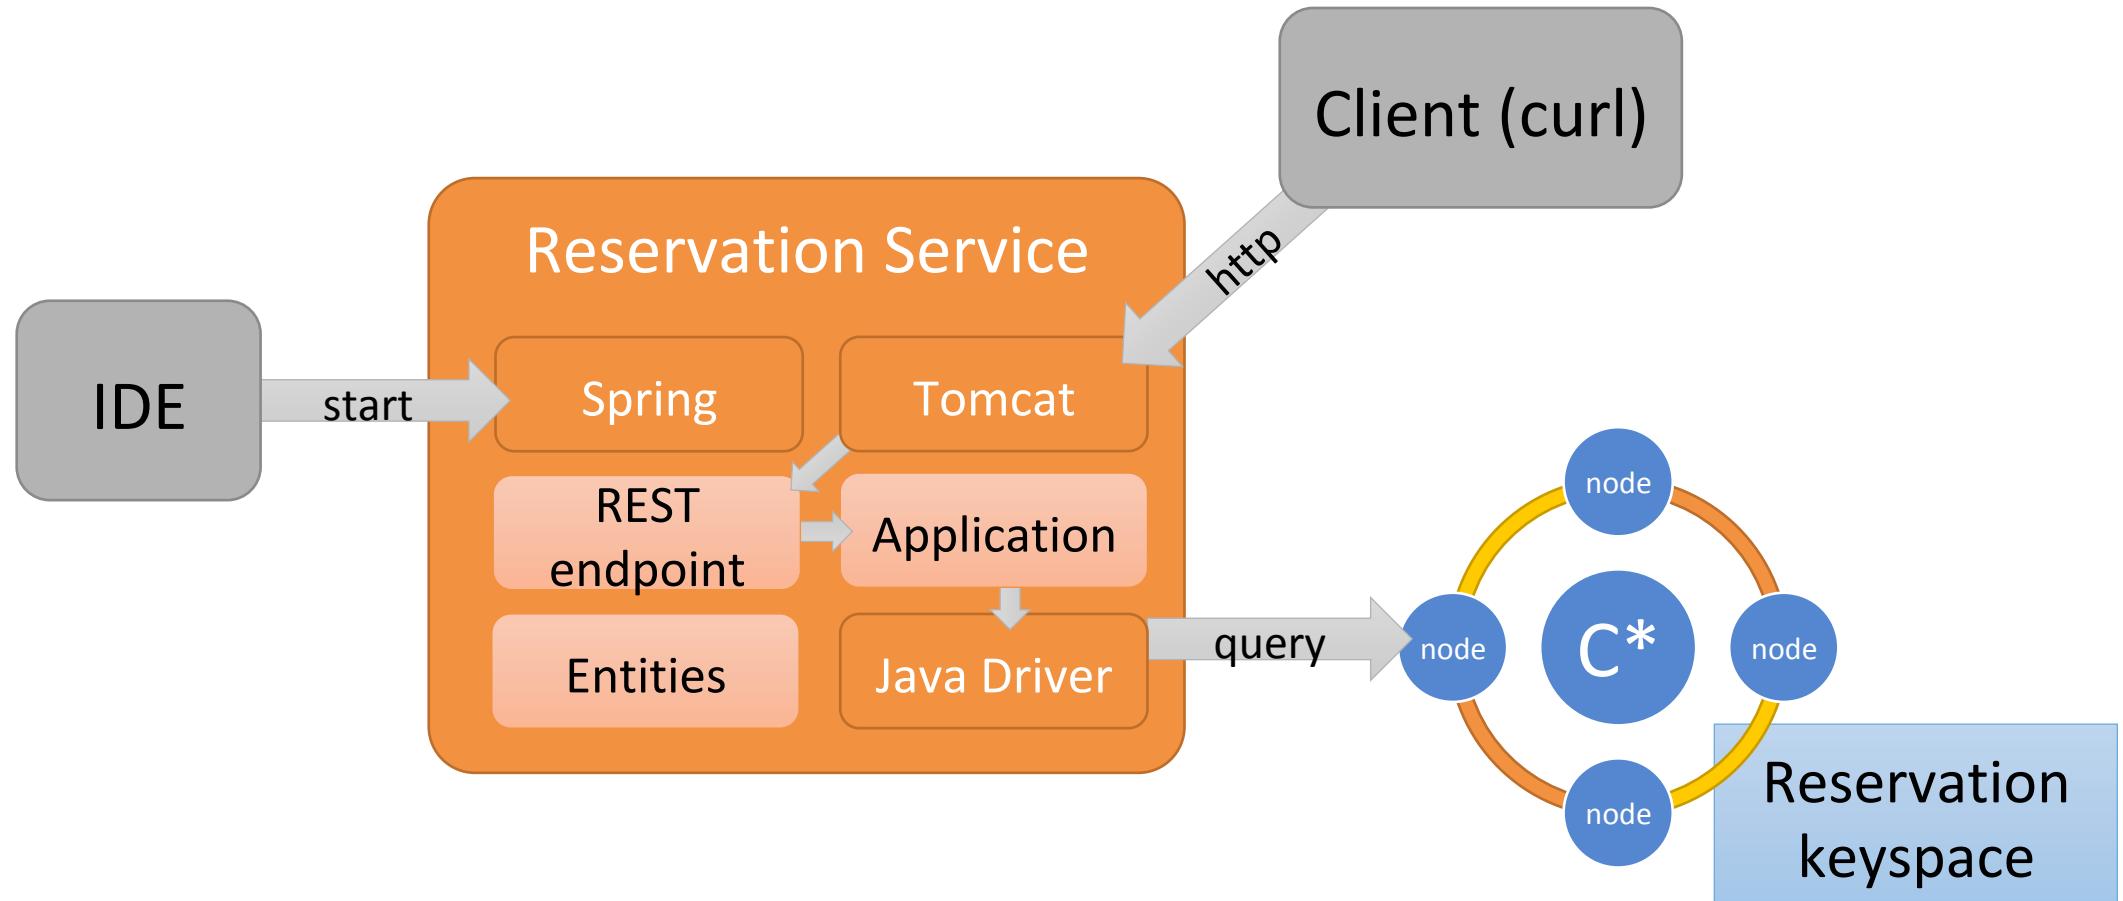
# DataStax Cassandra Drivers

- Common features:
  - Connection management
  - Creating and executing statements, and accessing the results
  - Synchronous and asynchronous execution
  - Object mapping
  - Logging and metrics
  - Policy management
  - Threading, networking and resource management
  - Schema access / management

# Our Project: Reservation Service

- Goals
  - Create a microservice implementation that persists data using Cassandra
  - Leverage best practices
  - Complete application except for the Cassandra bits
- Choices
  - Java
  - IntelliJ IDEA Community Edition
  - Spring Boot
  - RESTful API

# Reservation Service Architecture



# Exercise 3a – Reservation Service Setup

- Goal:
  - Set up the reservation service in our dev environment
- Steps
  - The sample application code has already been cloned from GitHub

```
$ git clone git://github.com/jeffreyscarpenter/reservation-service.git  
-or-  
$ cd reservation-service; git pull
```

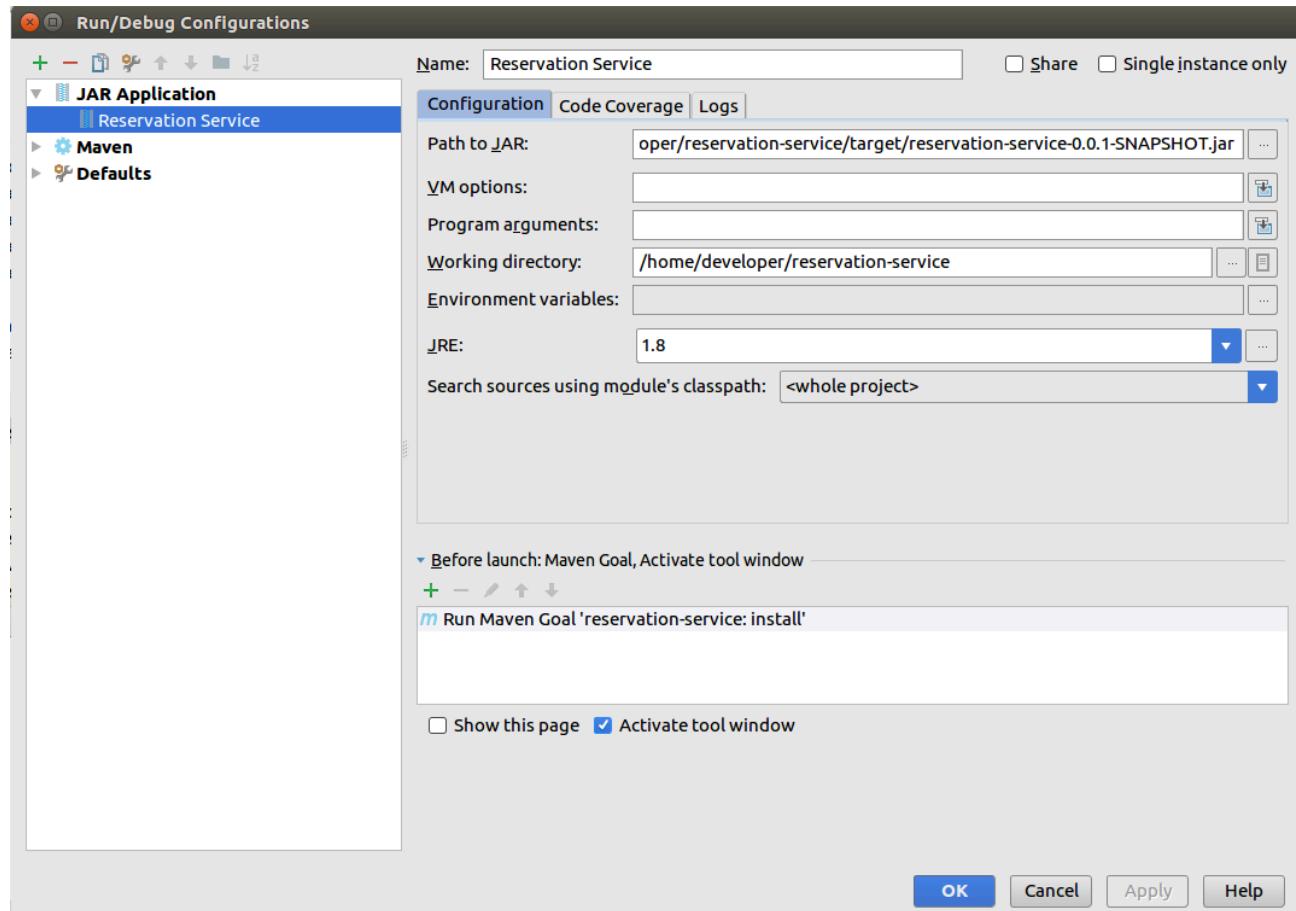
- Run *cqlsh* and create schema

```
cqlsh> SOURCE '~/reservation-service/src/main/resources/reservation.cql'
```

# Exercise 3a – Reservation Service Setup

- Instructions
  - Open IntelliJ, and select “Reservation Service” from recent projects
    - For reference: the project was created by “Select to create a new project from existing source” and configuring the project as a Maven project on import
  - Run the application
    - IntelliJ Ultimate: Run -> Run ReservationServiceApplication
    - IntelliJ Community Edition: Run -> Run Reservation Service
      - custom run config – see next slide
  - Eclipse is available on the VM but I recommend using IntelliJ
    - That’s how I set up the instructions
    - Better syntax highlighting
    - Better Spring Boot integration

# Exercise 3a – Reservation Service Setup



(This is just FYI, I've created this run configuration for you since we're using IntelliJ Community Edition, you don't have to do this in Ultimate)

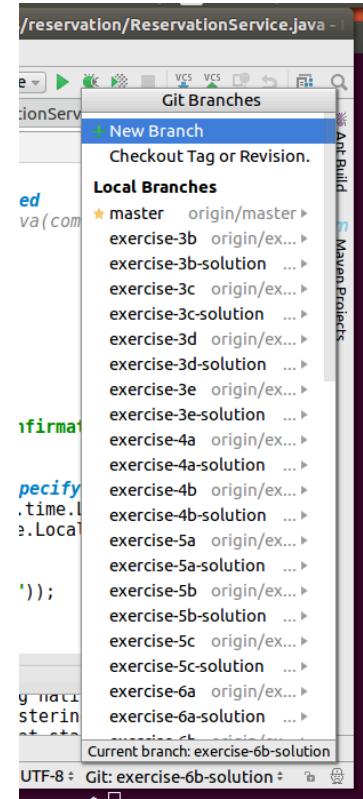
# Exercise 3a – Reservation Service Setup

- Examining the code:
  - Note the contents of the following files and directories, and then we'll discuss:
    - README.md
    - pom.mvn
    - /src
    - /resources
  - See if you can find the answers to these questions:
    - How do you access code samples for the various exercises?
    - What CQL type do we use to represent the start date and time of the reservation?
    - What Java type do the start date and time of the reservation get mapped to?

# Exercise 3a – Reservation Service Setup

- Questions and Answers

- How do you access code samples for the various exercises?
  - The start and end point for each exercise are tagged in Git
    - git checkout exercise-3b
    - git checkout exercise-3b-solution
  - You can also change branch you are looking at in IntelliJ IDEA (see picture)
- What CQL type do we use to represent the start date and time of the reservation?
  - date (not timestamp)
- What Java type do the start date and time of the reservation get mapped to?
  - java.time.LocalDate – I've chosen to use this instead of java.util.Date
  - This will be easier in some ways, harder in others, will use codecs later to simplify



# Reservation Service Contents

- `ReservationServiceApplication.java`
  - Spring boot auto-generated main function
- `ReservationServiceRestController.java`
  - Simple RESTful endpoint that turns HTTP requests into calls on `ReservationService`
- `ReservationService.java`
  - Contains the data storage logic of the service – main focus of our exercises
- `Reservation.java`
  - Entity definition
- `ConfirmationNumberGenerator.java`
  - Simple function to generate 6-digit alphanumeric confirmation numbers
  - The service is responsible for generating the confirmation numbers, not clients

# Reservation Service Caveats

- Not production-ready
  - Minimal data validation
  - Minimal fault handling
  - Schema makes use of Strings as identifiers instead of UUIDs in some cases
  - All code in single package (com.cassandraguide.services.reservation)
- These allow us to focus on the Cassandra bits
  - Checkout KillrVideo for a more robust reference application

# DataStax Drivers - Key Concepts

Class	Description
Cluster & ClusterBuilder	Representation of a cluster and default configuration options
Session	Manages connection to a cluster and associated resources
Statement	Represents a CQL query that can be executed: SimpleStatement, PreparedStatement -> BoundStatement, BuiltStatement
ResultSet	Result of executing a Statement, including metadata and rows matching select criteria
QueryBuilder	Fluent API for creating queries (BuiltStatements)
Mapper (Java, C#, Python)	Template for mapping Java classes to tables
*Policy (Various classes)	Multiple classes for configuring performance of the driver
*Codec (Various classes)	Manage conversions between Java and CQL types

\* Topics covered in Part 2 of this course

# DataStax Java Driver – Maven Configuration

- Update Maven pom.xml file for your project to include the Java Driver:

```
<dependency>
  <groupId>com.datastax.cassandra</groupId>
  <artifactId>cassandra-driver-core</artifactId>
  <version>3.6.0</version>
</dependency>
```

- (IntelliJ) Right click on project, select “Maven->Reimport”
  - You may also see an option to “auto import”

<http://docs.datastax.com/en/developer/java-driver/3.6/index.html> (“Getting the Driver”)

# DataStax Java Driver – Cluster and ClusterBuilder

- ClusterBuilder is a fluent interface for creating a cluster including:
  - Contact points (InetAddress)
  - Policies (we'll cover later)

```
Cluster cluster = Cluster.builder()  
    .addContactPoints("127.0.0.1")  
    .withLoadBalancingPolicy( ... )  
    .withRetryPolicy( ... )  
    .withReconnectionPolicy( ... )  
    .withPoolingOptions( ... )  
    .withQueryOptions( ... )  
    .build();
```

# DataStax Java Driver - Session

- Session represents a connection to the cluster with allocated resources
  - Thread pools
  - TCP connections
- Relatively heavyweight object
  - Create and reuse across multiple calls
- Can optionally bind to a keyspace

```
// Session bound to a keyspace
Session session =
    cluster.connect("reservation");

// Session not bound to a keyspace
Session anotherSession =
    cluster.connect();
```

# DataStax Java Driver – SimpleStatement

- Use Session to execute statements
- Can execute raw CQL as a string, or create a SimpleStatement and pass to Session.execute()
- Tip: instead of concatenating strings, use parameterization

```
session.execute("SELECT * FROM reservation"); // just a raw string

SimpleStatement statement = new SimpleStatement("SELECT * FROM reservation");
session.executeAsync(statement);

statement = new SimpleStatement("SELECT * FROM reservation WHERE confirmation_number = ?",
confirmationNumber); // parameterization
```

# Exercise 3b – DataStax Java Driver Setup and Simple Statement

- Goals:

- Configure project dependencies and connect to node on local machine
- Implement TODOs in ReservationService methods using SimpleStatement

- Steps

```
reservation-service$ git checkout exercise-3b
```

- Configure maven dependencies for our project in pom.xml
- Configure the ReservationService class to create a Cluster to connect to our local Cassandra node
  - Specify localhost as the connection point (“127.0.0.1”)
- Use the Cluster to create a Session bound to the “reservation” keyspace

# Exercise 3b - DataStax Java Driver Setup and Simple Statement

- Methods to update:
  - `createReservation()`
  - `retrieveReservation()`
  - `updateReservation()`
  - `getAllReservations()`
  - `deleteReservation()`
- Notes
  - Don't worry about `searchReservationsByHotelDate()` in any of today's exercises
  - Try to get at least the first two done, we can do the rest together

# Exercise 3b - DataStax Java Driver Setup and Simple Statement

- Test:
  - Use the *curl* command to send simple HTTP requests to the service
  - This file contains sample *curl* commands for simple testing:

```
$ more ~/reservation-service/src/test/simple_test_script.sh
```

- Take ID returned from create (POST) and substitute that in retrieve (GET)

# Testing with *curl*

```
# createReservation() (The “-i” option prints out headers, including location for GET)
curl -i -X POST --data '{"hotelId": "NY456", "startDate": "2017-06-08", "endDate": "2017-06-10", "roomNumber": "111", "guestId": "1b4d86f4-ccff-4256-a63d-45c905df2677"}' -H "Content-Type: application/json" http://localhost:8080/reservations/
# retrieveReservation()
curl http://localhost:8080/reservations/RS2G0Z
# updateReservation()
curl -X PUT --data '{"confirmationNumber": "RS2G0Z", "hotelId": "AZ123", "startDate": "2017-06-08", "endDate": "2017-06-14", "roomNumber": "999", "guestId": "1b4d86f4-ccff-4256-a63d-45c905df2677"}' -H "Content-Type: application/json" http://localhost:8080/reservations/
RS2G0Z
# getAllReservations()
curl http://localhost:8080/reservations/
# deleteReservation()
curl -X DELETE http://localhost:8080/reservations/RS2G0Z
```

# Exercise 3b - DataStax Java Driver Setup and Simple Statement

- Finishing up:
  - To save your work (locally):

```
reservation-service$ git commit -a -m "my work"
```

- To see the solution

```
reservation-service$ git checkout exercise-3b-solution
```

- Discard your changes / revert (works on all exercises)

```
reservation-service$ git checkout -- .
```

# DataStax Java Driver – Prepared Statement

- Created on client, driver registers with each node in the cluster
- Keep a copy for use on incoming requests
- When executed on client, only the prepared statement ID and values are sent
- Pros: efficiency, security

```
import com.datastax.driver.core.PreparedStatement;  
...  
PreparedStatement hotelInsertPrepared = session.prepare(  
    "INSERT INTO hotels (id, name, phone) VALUES (?, ?, ?)");
```

# DataStax Java Driver – Bound Statement

- To execute a query based on a PreparedStatement, we provide values to create a BoundStatement
- The BoundStatement can then be executed using the Session

```
// PreparedStatement listed attributes: id, name, phone  
BoundStatement hotelInsertBound = hotelInsertPrepared.bind(  
    id, "Super Hotel at WestWorld", "1-888-999-9999");  
  
ResultSet hotelInsertResult = session.execute(hotelInsertBound);
```

# Exercise 3c – Prepared Statement

- Goal:
  - Re-Implement the Reservation Service operations using PreparedStatement
- Steps:

```
reservation-service$ git checkout exercise-3c
```

- Add import statements
- Create PreparedStatement in constructor
- Create/Execute BoundStatements in each operation
  - createReservation()
  - retrieveReservation()
  - updateReservation()
  - getAllReservations()
  - deleteReservation()

# Exercise 3c – Prepared Statement

- Finishing up:
  - To save your work (locally):

```
reservation-service$ git commit -a -m "my work"
```

- To see the solution

```
reservation-service$ git checkout exercise-3c-solution
```

- Questions?

# DataStax Java Driver – Query Builder

- Fluent API for building up statements a portion at a time
- More programmatic, less string manipulation

```
import com.datastax.driver.core.querybuilder.BuiltStatement;
import com.datastax.driver.core.querybuilder.QueryBuilder;
import static com.datastax.driver.core.querybuilder.QueryBuilder.eq;

BuiltStatement hotelInsertBuilt = QueryBuilder.insertInto("hotels")
    .value("id", id)
    .value("name", "Super Hotel at WestWorld")
    .value("phone", "1-888-999-9999");

BuiltStatement hotelSelectBuilt = QueryBuilder.select().all().
    from("hotels").where(eq("id", id));
```

See more examples at: <https://gist.github.com/yangzhe1991/10349122>

# Exercise 3d – Query Builder

- Goal:
  - Re-implement the Reservation Service operations using QueryBuilder
- Steps:

```
reservation-service$ git checkout exercise-3d
```

- Add import statements
- Create a QueryBuilder we can reuse in each operation
- Create/execute BuiltStatements in each operation
  - createReservation()
  - retrieveReservation()
  - updateReservation()
  - getAllReservations()
  - deleteReservation()

# Exercise 3d – Query Builder

- Finishing up:
  - To save your work (locally):

```
reservation-service$ git commit -a -m "my work"
```

- To see the solution

```
reservation-service$ git checkout exercise-3d-solution
```

- Questions?

# DataStax Java Driver – Mapper

- Abstracts details of mapping Java attributes to/from CQL types and UDTs
- Works via annotations
- Packaged separately from the driver – pom.xml update required

```
<dependency>
  <groupId>com.datastax.cassandra</groupId>
  <artifactId>cassandra-driver-core</artifactId>
  <version>3.3.0</version>
</dependency>
```

# DataStax Java Driver – Mapper Annotations

- Data Type Class

```
import com.datastax.driver.mapping.annotations.Column;
import com.datastax.driver.mapping.annotations.PartitionKey;
import com.datastax.driver.mapping.annotations.Table;

@Table(keyspace = "hotel", name = "hotels") public class Hotel {
    @PartitionKey @Column(name = "id") private String id;
    @Column(name = "name") private String name;
    // Other attributes, Get/set>equals/hashCode methods
}
```

# DataStax Java Driver – Using the Mapper

- Executing queries

```
import com.datastax.driver.mapping.MappingManager;
import com.datastax.driver.mapping.Mapper;

MappingManager mappingManager = new MappingManager(session);
Mapper<Hotel> hotelMapper = mappingManager.mapper(Hotel.class);

hotelMapper.save(hotel); // pass a Hotel object to insert/update
Hotel hotel = hotelMapper.get(hotelId); // pass primary key to load
hotelMapper.delete(hotelId) // pass primary key to delete
```

# Exercise 3e – Mapper

- Goal:
  - Re-implement the Reservation Service operations using Mapper
- Steps:

```
reservation-service$ git checkout exercise-3e
```

- Update pom.xml to include mapper
- Complete the annotations in ReservationByConfirmation.java
- Imports and implement statements in ReservationService.java
  - createReservation()
  - retrieveReservation()
  - updateReservation()
  - getAllReservations()
  - deleteReservation()

# Exercise 3e - Mapper

- Finishing up:
  - To save your work (locally):

```
reservation-service$ git commit -a -m "my work"
```

- To see the solution
- Questions?

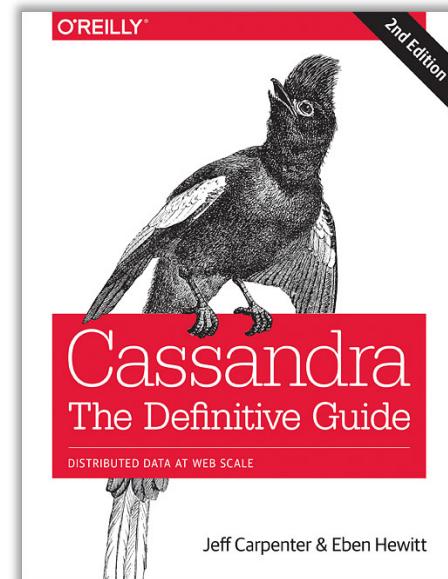
```
reservation-service$ git checkout exercise-3e-solution
```

# DataStax Java Driver – Statements

- So which is the best? It depends on your situation and preferences
  - PreparedStatement – Efficiency and safety for structured queries
  - SimpleStatement – Ad-hoc queries with variable structure
    - Always use parameterization for security
  - QueryBuilder – clean looking code
  - ObjectMapper – leveraging simple domain models, abstracting driver

# Preparing for Part 2

- Finish up any remaining work from activities
- Reinforce today's learning with readings in *Cassandra, The Definitive Guide*, 2nd Ed:
  - Chapter 3 – Installing Cassandra
  - Chapter 4 – The Cassandra Query Language
    - Especially “CQL Types”
- Read before the next session:
  - Chapter 6 – The Cassandra Architecture
  - Chapter 7 – Configuring Cassandra
- Make sure you keep your VM around...



# End of Part 1

See you in two days!